ResolveAI Data-to-UI Mapping

Where each field on screen comes from: database columns in Part A, and log, deploy and metric files in Part B.

Based on the code at commit 675082e on main, 2 November 2026. Demo dataset: all incidents, logs and metrics are synthetic.

15	138	79	30
database tables mapped	columns checked against the schema	columns displayed, directly or formatted	columns never displayed

The most important thing to know: the Logs tab does not read the log files. It reads the seeded `log_events` table. The files in `demo-systems/logs` appear only in the System checks tab and in cited evidence. The same split applies to deployments and metrics, so the two views can differ.

Route 1: database tables

Seed.cs — Fixed demo data → **SQLite** — log_events, deployments, metric_snapshots → **EvidenceLoader** — Picks by service and time → **Tabs** — Signals, Logs, Deployments

Available before any investigation runs. Stored log text is shown as is.

Route 2: files and SQL Server, through MCP tools

demo-systems/ — Logs, deploys, metrics; live SQL Server → **MCP tools** — Read-only checks → **tool_calls** — Summary and rows as JSON → **System checks** — Cards and tables

Only exists after an investigation. Results appear exactly as returned.

How to read the “On screen” column

Direct — Value appears as stored.

Formatted — Appears after formatting, such as a date, chip or JSON list.

Drives UI — Not printed, but it filters, orders or switches something on screen.

Via evidence — Stored JSON or input used elsewhere; reaches the screen through another field, or not at all.

Not shown — Internal, or stored but never displayed.

Part A: Database to UI

Every table and column in the SQLite database, in the order you meet them in the app. Tabs named here are the tabs on the incident page.

Queue and incident header

`incidents`

Seeded by Seed.cs

One row per active or resolved incident. Read by the queue and the incident header.

Column	On screen	UI location	UI field and notes
`id`	Not shown	Internal	Primary key. Links the other tables; the UI uses key instead.
`key`	Direct	Queue card · Home spotlight · Incident header	Grey ID such as INC-1042. Also the page URL, /incidents/INC-1042.
`title`	Direct	Queue card · Home spotlight · Incident header	Bold card title; large heading on the other two.
`service_name`	Direct	Queue card · Home spotlight · Incident header · Logs tab intro	Coloured service chip. Also fills the All services filter.
`severity`	Formatted	Queue card · Home spotlight · Incident header	SEV chip and the coloured left edge of the card. Also the severity filter and SEV-1 tile.
`status`	Formatted	Queue card · Incident header	Status chip. Also the status filter; resolved incidents sort last and leave the Active tile.
`environment`	Direct	Incident header	Small grey text after the service chip, such as production.
`started_at`	Formatted	Queue card · Incident header · Deployments tab	Started time and Age on the card; “Started …” in the header; “N min before start” on deployments.
`detected_at`	Formatted	Incident header	“detected …” after the start time.
`resolved_at`	Formatted	Incident header	“resolved …”, only when it has a value.
`description`	Direct	Incident header	Paragraph under the title. Hidden while the header is collapsed after a run.
`impact_summary`	Direct	Queue card · Home spotlight · Incident header	One-line summary on the card; red “Impact:” box in the header.
`created_at`	Not shown	Internal	Never read by the app.

Evidence tabs

`log_events`

Seeded by Seed.cs

The seeded log lines behind the Logs tab. This table is not filled from the log files.

Column	On screen	UI location	UI field and notes
`id`	Not shown	Internal	Forms the evidence reference log:<id> used to jump to the row.
`incident_id`	Drives UI	Logs tab (filter)	A row appears if it matches the incident’s service or is linked by this column.
`timestamp`	Formatted	Logs tab: Time · Signals tab · AI findings	Shown as 10:04:12 UTC. Also First seen / Last seen, and “First error at …” in Observed facts.
`service_name`	Direct	Logs tab: Service	Service column of the log table.

Column	On screen	UI location	UI field and notes
level	Formatted	Logs tab: Level	Coloured INFO / WARN / ERROR badge. Observed facts looks for the first ERROR.
message	Direct	Logs tab: Message · AI findings	Stored text, not redacted. Also quoted in “First error at …”.
error_signature	Direct	Logs tab · Signals tab: Error signals	Grey tag after the message, and the Signature column of Error signals.
trace_id	Not shown	Internal	Read by the query but never displayed. NULL in the demo data.
source	Via evidence	Evidence packet only	Becomes the packet source, such as app/payment-api, but is never displayed.

deployments

Seeded by Seed.cs

Seeded release history behind the Deployments tab.

Column	On screen	UI location	UI field and notes
id	Not shown	Internal	Forms the evidence reference dep:<id>.
service_name	Direct	Deployments tab	First half of the badge, such as “payment-api v2.8.1”. Indigo when it is the incident’s service.
environment	Direct	Deployments tab	After the time: “… · production”.
version	Direct	Deployments tab	Second half of the badge.
deployed_at	Formatted	Deployments tab · AI findings	Date and time, plus “N min before start”. Observed facts lists releases up to 2 hours before.
change_summary	Direct	Deployments tab · AI findings	Main text of the row.
commit_ref	Direct	Deployments tab	“commit 7d03b1f”, when present.
incident_id	Not shown	Internal	Not used. Deployments are chosen by service and time window.

metric_snapshots

Seeded by Seed.cs

One baseline and peak per metric. These are not the time series in the metrics JSON files.

Column	On screen	UI location	UI field and notes
id	Not shown	Internal	Forms the evidence reference met:<id>.
incident_id	Drives UI	Signals tab (filter)	Selects the incident’s metrics.
name	Direct	Signals tab: Metrics card	Small label at the top of the metric card.
unit	Direct	Signals tab: Metrics card	Small text after the peak, such as %, ms or connections.
baseline	Formatted	Signals tab: Metrics card	Number before the arrow. Also used for the ratio chip.
peak	Formatted	Signals tab: Metrics card	Large number after the arrow; red if it moved sharply.
observed_at	Not shown	Internal	Read by the query but never displayed.

`timeline_events`

Seeded by Seed.cs

Seeded incident timeline.

Column	On screen	UI location	UI field and notes
`id`	Not shown	Internal	Forms the evidence reference tl:<id>.
`incident_id`	Drives UI	Signals tab (filter)	Selects the incident's events.
`timestamp`	Formatted	Signals tab: Timeline	"10:05:00 UTC · alert".
`kind`	Formatted	Signals tab: Timeline	Text after the time. Also sets the dot colour.
`description`	Direct	Signals tab: Timeline	Main text of the event.

`knowledge_documents`

Seeded by Seed.cs

Runbooks, guides and historical RCA notes.

Column	On screen	UI location	UI field and notes
`id`	Not shown	Internal	Forms the evidence reference doc:<id>; loads a historical RCA.
`title`	Direct	Signals tab: Runbooks & guides · Similar tab	Bold line above the body.
`doc_type`	Drives UI	Signals tab (filter)	Only runbook and guide appear in Runbooks & guides. rca documents open from a similar incident.
`service_name`	Drives UI	Signals tab (filter)	Same service, or NULL for general guides.
`body`	Direct	Signals tab: Runbooks & guides · Similar tab	Runbook text; text of "View historical RCA".
`source_uri`	Not shown	Internal	Never read by the UI. NULL in the demo data.
`tags`	Not shown	Internal	Only copied into knowledge_fts, which nothing searches.
`created_at`	Not shown	Internal	Never read by the UI.
`updated_at`	Not shown	Internal	Never read by the UI.

AI findings and similar incidents

`investigations`

Written by the investigation pipeline

One row per investigation run. A new run adds a row; earlier runs stay.

Column	On screen	UI location	UI field and notes
`id`	Not shown	Internal	Value behind the run selector.
`incident_id`	Drives UI	Incident page (filter)	Finds the runs for the incident.
`status`	Drives UI	AI findings · Run selector	Chooses the progress board, the error box or the results. Also in the selector text.
`stage`	Direct	AI findings: progress board	Line under "Investigation in progress". Also drives the step list.
`question`	Direct	AI findings: Inferred analysis	"Question asked: …".

Column	On screen	UI location	UI field and notes
`llm_mode`	Formatted	AI findings · Run selector	Badge wording: “AI-generated (…)” or “Recorded analysis”.
`evidence_packet`	Via evidence	AI findings	Stored JSON. Supplies each evidence excerpt and the “N evidence records” count.
`model_name`	Direct	AI findings: AI badge	Inside “AI-generated (claude-opus-4-8) — not confirmed”.
`prompt_version`	Direct	AI findings: Inferred analysis	“prompt inv-2026-10-28.1”.
`summary`	Direct	AI findings: Inferred analysis	Summary paragraph.
`insufficient_evidence`	Drives UI	AI findings	Shows the amber “Insufficient evidence” box instead of the leading hypothesis.
`limitations`	Direct	AI findings: Limitations	JSON list, shown as bullets.
`validation_issues`	Direct	AI findings: Validation notes	JSON list, shown as bullets.
`usage`	Not shown	Internal	Written, never read. Token counts on screen come from trace.
`started_at`	Formatted	Run selector	Date and time in each run’s option.
`completed_at`	Formatted	AI findings: Inferred analysis	“completed …”.
`error_message`	Direct	AI findings: error box	Shown when a run failed.
`trace`	Direct	AI findings: How this was produced	JSON list of stages, durations and details, including token counts.

hypotheses

Written by the investigation pipeline

Ranked causes from one run.

Column	On screen	UI location	UI field and notes
`id`	Not shown	Internal	Links votes and evidence.
`investigation_id`	Drives UI	AI findings (filter)	Selects the run’s hypotheses.
`rank`	Direct	AI findings: hypothesis card	Number in the dark circle. Rank 1 is also the Leading hypothesis.
`title`	Direct	AI findings · Run selector · Verified form	Card and leading-hypothesis heading. Shortened in the selector and the form’s list.
`explanation`	Direct	AI findings: hypothesis card	Paragraph under the heading.
`confidence`	Formatted	AI findings	Low / medium / high chip.
`score`	Not shown	Internal	Stored and loaded, never displayed.
`speculative`	Drives UI	AI findings	Amber “speculative” tag and warning.
`counter_evidence`	Direct	AI findings: Counter-evidence	JSON list, shown as bullets.
`unknowns`	Direct	AI findings: Unknowns	JSON list, shown as bullets.
`recommended_checks`	Direct	AI findings: Recommended checks · Check this first	JSON list: Check, Why, Expected signal. The first appears in “Check this first”.

evidence_links

Written by the investigation pipeline

Which records each hypothesis cites.

Column	On screen	UI location	UI field and notes
hypothesis_id	Drives UI	AI findings (filter)	Selects a hypothesis's citations.
evidence_ref	Via evidence	AI findings: What backs it · Supporting evidence	Looked up in evidence_packet for the chip and excerpt. Click to jump to the source.
position	Drives UI	AI findings (order)	Citation order. "What backs it" shows the first 3.

historical_incidents

Seeded by Seed.cs

Past incidents searched for similar symptoms.

Column	On screen	UI location	UI field and notes
id	Not shown	Internal	Forms the evidence reference hist:<id>.
key	Direct	Similar tab	Grey ID on the card, such as INC-0881.
title	Direct	Similar tab	Bold card heading.
service_name	Direct	Similar tab	Grey text after the title.
symptoms	Via evidence	AI findings, when cited	Not on the card. Searched through historical_fts, and part of the cited evidence text.
root_cause	Direct	Similar tab: Prior root cause	Also part of the cited evidence text.
resolution	Direct	Similar tab: Resolution	Also part of the cited evidence text.
rca_document_id	Drives UI	Similar tab	Shows the "View historical RCA" button when set.
resolved_at	Formatted	Similar tab	"resolved 2026-03-14", date only.
tags	Not shown	Internal	Only copied into historical_fts, where it helps matching.

investigation_similar_incidents

Written by the investigation pipeline

Past incidents matched by one run.

Column	On screen	UI location	UI field and notes
investigation_id	Drives UI	Similar tab (filter)	Selects the run's matches.
historical_incident_id	Drives UI	Similar tab (join)	Joins to historical_incidents.
rank	Drives UI	Similar tab (order)	Card order.
match_score	Not shown	Internal	Loaded, never displayed.
matched_terms	Direct	Similar tab	Grey term tags; first 8 shown.
similarity_reason	Direct	Similar tab: Why it matches	Sentence on the card.

RCA, feedback and system checks

`rca_drafts`

Written by RCA generation and the RCA editor

Each generated RCA version, including a person's edits.

Column	On screen	UI location	UI field and notes
`id`	Not shown	Internal	Value behind the version selector; used in the export link.
`incident_id`	Drives UI	RCA tab (filter)	Selects the incident's drafts.
`investigation_id`	Not shown	Internal	Loaded, never displayed.
`version`	Direct	RCA tab · tab label	"v1" on the tab and editor; "Version 1" in the selector.
`status`	Formatted	RCA tab	draft / reviewed badge. Latest reviewed draft helps set Verified on the queue.
`impact`	Direct	RCA tab: Impact	Editable text box.
`timeline`	Direct	RCA tab: Timeline (UTC)	JSON list, read-only, with "source" buttons.
`root_cause`	Direct	RCA tab: Root cause	Editable text box.
`contributing_factors`	Direct	RCA tab: Contributing factors	JSON list, edited one per line.
`resolution_steps`	Direct	RCA tab: Resolution	JSON list, edited one per line.
`prevention_actions`	Direct	RCA tab: Preventive actions	JSON list, edited one per line.
`evidence_ids`	Not shown	Internal	Loaded, never displayed. The timeline's own references drive the source buttons.
`generated_at`	Formatted	RCA tab	"generated …".
`updated_at`	Formatted	RCA tab	"edited …".
`validation_issues`	Direct	RCA tab: Check before publishing	JSON list, shown as bullets.
`llm_mode`	Formatted	RCA tab	"AI-generated draft" or "Recorded draft".

`feedback`

Written from the AI findings tab

Votes and what a person verified.

Column	On screen	UI location	UI field and notes
`id`	Not shown	Internal	Primary key.
`investigation_id`	Drives UI	Queue · RCA generation	Ties feedback to a run.
`hypothesis_id`	Drives UI	AI findings	Ties a vote to a card. Set by "Related hypothesis" in the Verified form.
`helpful`	Formatted	AI findings: Useful / Not useful	Latest vote highlights the button.
`comment`	Via evidence	Input: Investigation notes	Not shown again. Passed to the next RCA draft as an engineer note.
`verified_root_cause`	Drives UI	Input: Verified root cause	Not shown again. Used by the next RCA draft, and helps set Verified on the queue.
`created_at`	Drives UI	AI findings (order)	Finds the latest vote.

`tool_calls`

Written by the investigation pipeline

Each read-only system check in a run. This is how log files reach the UI.

Column	On screen	UI location	UI field and notes
`id`	Not shown	Internal	Primary key. The on-screen ID comes from seq.
`investigation_id`	Drives UI	System checks (filter)	Selects the run's checks.
`seq`	Formatted	System checks	ID chip on the card, such as mcp:1. Findings cite this.
`phase`	Formatted	System checks · Trace	Baseline or Follow-up badge.
`system`	Formatted	System checks	Group heading and colour: Database, Files & logs, Deploys & Git, Metrics & host.
`server`	Not shown	Internal	Loaded, never displayed on this tab.
`tool`	Direct	System checks · Trace	Coloured chip, such as files_search.
`args`	Formatted	System checks	name=value chips under the title.
`ok`	Drives UI	System checks · Trace	"failed" badge and red summary when false.
`summary`	Direct	System checks · Trace	One-line result. Also the excerpt on a citation.
`rows`	Formatted	System checks	JSON rows, as a table or raw JSON. See Part B.
`truncated`	Drives UI	System checks	"truncated" badge.
`duration_ms`	Direct	System checks · Trace	"N ms".
`started_at`	Not shown	Internal	Loaded, never displayed.

`_migrations`

Written at startup

Records which schema migrations have run.

Column	On screen	UI location	UI field and notes
`name`	Not shown	Internal	Never shown.
`applied_at`	Not shown	Internal	Never shown.

`knowledge_fts` and `historical_fts` are search indexes, not displayed tables. `historical_fts` powers the similar-incident search. `knowledge_fts` is filled when the data is seeded, but no code searches it.

Fields with no column of their own

UI field	Location	Worked out from
Review state: Not investigated / Needs review / Verified	Queue card, tiles, filter	SQL in Incidents.cs. `investigations.status`, `feedback.verified_root_cause` and the latest `rca_drafts.status`.
Active, SEV-1, Needs review, Verified tiles	Home page	Counts of `incidents` using status, severity and review state.
Age	Queue card	Current time minus `incidents.started_at`.

UI field	Location	Worked out from
Error signals: Count, First seen, Last seen	Signals tab	Grouped from the loaded `log_events` rows by `error_signature`.
Observed facts	AI findings	Sentences built from `log_events`, `deployments` and `metric_snapshots`. No model involved.
Ratio chip (×140, ≈ flat, new)	Metrics card	`peak` ÷ `baseline`.
"N min before start"	Deployments tab	`incidents.started_at` minus `deployments.deployed_at`.
Logs tab count	Tab label	Rows loaded from `log_events`, up to 200.
N evidence records	AI findings	Items in `investigations.evidence_packet`.
N of M left today	Incident header	Today's live `investigations` and `rca_drafts` rows, against `LIVE_DAILY_LIMIT`.

Fields you type, and where they are saved

UI input	Location	Saved to
Optional question	Incident header	`investigations.question`
Run investigation	Incident header	New `investigations` row. The pipeline fills the rest.
Useful / Not useful	Hypothesis card	`feedback.helpful`, `feedback.hypothesis_id`
Verified root cause	Record what was verified	`feedback.verified_root_cause`
Investigation notes	Record what was verified	`feedback.comment`
Related hypothesis	Record what was verified	`feedback.hypothesis_id`
Impact, Root cause	RCA editor	`rca_drafts.impact`, `rca_drafts.root_cause`
Contributing factors, Resolution, Preventive actions	RCA editor	`contributing_factors`, `resolution_steps`, `prevention_actions`, one JSON item per line
Save & mark reviewed	RCA editor	`rca_drafts.status` = reviewed, and `updated_at`

Fields that do not come from the database

UI field	Location	Source
Data source: Demo dataset	Top bar	Fixed label.
LLM: live / recorded	Top bar	`.env`: `LLM_MODE`, `ANTHROPIC_API_KEY`, `LIVE_ACCESS_CODE`, `CLAUDE_MODEL`.
Connected / Not working / Turned off, tools	Systems page	Live status of the MCP servers at run time, not `tool_calls`.
Systems checked strip	Incident header	Live status, plus this run's `tool_calls`.

Part B: Logs and files to UI

One event through both routes: the payment-api timeout on INC-1042, from the real files and seed data.

Where	What it holds
Log file, line 6	`2026-10-27T10:04:12.403Z ERROR payment-api database connection acquisition timed out after 5000ms; request_id=req-781`
System checks row	`file` logs/payment-api/app.log · `line` 6 · `time` 2026-10-27T10:04:12.403Z · `text` the whole line above
log_events row	`timestamp` 2026-10-27T10:04:12Z · `level` ERROR · `service_name` payment-api · `message` payment-api: database connection acquisition timed out; request_id=req-781 · `error_signature` db_conn_acquire_timeout · `source` app
Logs tab row	Time 10:04:12 UTC · Level ERROR · Service payment-api · Message, then a grey db_conn_acquire_timeout tag
Cited in AI findings	Log ×3 chip · ERROR payment-api: database connection acquisition timed out; request_id=req-781 [signature=db_conn_acquire_timeout]

The wording and the timestamp precision differ, because the file and the table are separate data. Only the table has an error signature, so the Error signals card has nothing to do with the files.

Logs tab: `log_events` to table columns

UI column	Column	Rule
Time	`timestamp`	Shown as HH:MM:SS UTC. Rows are sorted by this column, up to 200.
Level	`level`	Coloured badge.
Service	`service_name`	Incident's service, plus other services' rows linked by incident_id, such as payments-db.
Message	`message` + `error_signature`	Message, then a grey tag when a signature exists.
(window)	`incidents.started_at`, `detected_at`	30 minutes before the start, to 1 hour after the later of start and detection.

System checks: log file lines to row columns

`files_search` scans the service's folder for each phrase and returns one row per matching line. `RowsTable` shows up to 7 columns and the first 8 rows, with the rest in "Show all" and "Show raw JSON". Headings are the field names in capitals with underscores removed.

File line	Row field	UI column	Notes
Path of the file	`file`	FILE	`logs/payment-api/app.log`
Line number	`line`	LINE	Counted while reading the file.
Leading timestamp	`time`	TIME	Pulled out by a pattern, and also used for the time window.
The whole line	`text`	TEXT	Level, service and message stay together. They are not separate columns.

Which files reach the UI, and how

File	Reaches the UI through	Note
`logs/<service>/app.log`	`files_search`, for ERROR and WARN, on every run	Lines must fall in the incident window.

File	Reaches the UI through	Note
`logs/payment-api/db-client.yaml`	`files_tail` or `files_check`, as a follow-up	No leading timestamps, so the windowed baseline search skips its lines.
`logs/document-service/storage.yaml`, `mounts.txt`	`files_tail` or `files_check`, as a follow-up	Same reason.
`deploys.json`	`deploys_recent` on every run; `deploys_change_summary` as a follow-up	Separate from the deployments table.
`metrics/<service>.json`	`metrics_series` on every run, for that service's key metrics	`search-api` and `checkout-web` have no metrics file.

Services and folders come from `mcp.config.json`, or from `mcp.config.example.json` when there isn't one. This table follows the example file.

Deploy and metric files to row columns

Source field	Tool	Row field / UI column	Notes
`deploys.json` `service`, `version`, `commit`, `author`	`deploys_recent`	service, version, commit, author	Same names. Row also has `kind` = deployment.
`deployed_at`	`deploys_recent`	time	Renamed. Also filters the window.
`change_summary`	`deploys_recent`	summary	Renamed and clipped.
`environment`	`deploys_recent`	environment	Defaults to production when absent.
`files_changed[]` `path`, `added`, `removed`	`deploys_change_summary`	path, added, removed	One row per changed file.
`patch`	`deploys_change_summary`	patch	One more row, shown last.
`metrics/<service>.json` points [time, value]	`metrics_series`	time, value	One row per point in the window.
`unit`	`metrics_series`	unit	Repeated on each row.
Baseline, peak, average	`metrics_series`	Summary sentence	Baseline is the first value in the window; peak and average are computed.

Each tool result on the System checks card

Card field	Column in tool_calls	Example
Tool chip	`tool`	files_search
Baseline / Follow-up	`phase`	Baseline runs every time; Follow-up is chosen by the model, or replayed in recorded mode
ID, row count, time	`seq`, `rows`, `duration_ms`	mcp:1 · 6 rows · 5 ms
Argument chips	`args`	root=payment-api, phrase=ERROR, from=..., to=..., max_matches=30
Summary sentence	`summary`	6 line(s) containing "ERROR" in logs/payment-api (2 file(s) scanned).
Result table	`rows`	Columns depend on the tool. See the tables above.

SQL Server checks

These query the live server, not a file, and only when SQL Server is enabled and the service has a database in the MCP config. On every run, that service gets `db_health`, `db_blocking`, `db_errors` and `db_deadlocks`. The others are available as follow-ups.

Tool	Row fields shown in System checks
db_health	section, then server (cpu_count, physical_memory_mb, committed_mb, target_mb, sqlserver_start_time, version, edition), sessions (database_name, sessions, active_requests, blocked_requests) or top_wait_since_start (wait_type, waiting_tasks_count, wait_time_ms)
db_blocking	role, then head_blocker (session_id, status, open_transaction_count, last_request_end_time, program_name, last_query_text) or blocked (session_id, blocking_session_id, database_name, wait_type, wait_ms, wait_resource, status, command, query_text, program_name)
db_errors	time, error_number, severity, message, captured_by, and database_name for resolveai_errors rows
db_deadlocks	time, process, spid, database_name, wait_resource, lock_mode, query_text, victim
db_long_queries	session_id, database_name, status, command, elapsed_s, cpu_ms, logical_reads, wait_type, query_text, program_name
db_space	section, then database (database_name, state_desc, recovery_model_desc, log_reuse_wait_desc, is_read_only) or file (file_name, type_desc, size_mb, used_mb, max_mb, growth)
db_failed_jobs	job_name, step_id, step_name, run_at, message

Where clicking a citation goes

Reference	Opens tab	Highlights	Source
log:<id>	Logs	The log row	log_events
dep:<id>	Deployments	The deployment	deployments
met:<id>, tl:<id>, doc:<id>	Signals & timeline	The metric, event or runbook	metric_snapshots, timeline_events, knowledge_documents
hist:<id>	Similar incidents	The similar-incident card	historical_incidents
mcp:<n>	System checks	The tool card	tool_calls

What is and isn't redacted

- **Not redacted:** the Logs tab and System checks show text as stored or returned. The demo data contains a deliberately fake token and password, so you will see them there.
- **Redacted:** evidence sent to the model, and the excerpts cited in AI findings, such as `Bearer [REDACTED]`.